Buzz: next step report

BongoHive AI Lab, 21 October 2026

What the next step was

PLAN section 11 lists three entry criteria for Phase 1. The third is a read-only inspection of each server (lsb_release, nvidia-smi, lspci, df -h, ss -tulpn, docker version) with the output saved in docs/inspection/. No servers exist yet, so I built the tooling now. When the temporary cloud VMs are ready, one command inspects each host and a second command says whether it fits the plan.

What was built

File	Purpose
scripts/inspect-server.sh	Runs the six commands for role "inference" or "app". Changes nothing. A missing tool is recorded as NOT-INSTALLED instead of aborting. Can run over ssh with no copy step.
scripts/check_inspection.py	Reads the output and gives OK / WARN / FAIL for Ubuntu 22.04 or 24.04, GPU size (picks the 16, 24 or 48 GB preset from PLAN section 4), driver R580 or newer, free disk, clashes with our ports, and Docker. Exits 1 on any FAIL.
tests/inspection/	4 tests: good host passes, port clash fails, missing GPU fails, 8 GB GPU fails.
docs/inspection/	Folder where each host's output will be saved.
docs/CHANGELOG.md	One line added under Added.

How to use it

```
ssh user@gpu-vm 'bash -s' -- inference < scripts/inspect-server.sh \
  | tee docs/inspection/gpu-vm.txt
python3 scripts/check_inspection.py docs/inspection/gpu-vm.txt
```

What I ran

The 4 new tests pass. I ran the script on this Mac as a dry run. It completed and recorded NOT-INSTALLED for lsb_release, nvidia-smi and ss. The checker then reported FAIL for Ubuntu, GPU memory and driver, and WARN for disk and Docker, which is correct because a Mac is not an inference server. I deleted that output so it can't be mistaken for a real inspection.

```
[FAIL] Ubuntu: lsb_release missing or unreadable
[FAIL] GPU memory: no GPU reported by nvidia-smi
[FAIL] NVIDIA driver: not found
[WARN] Disk on /: could not read the / line
[OK  ] Ports free: none of ours are taken
[WARN] Docker: not installed yet
```

Not done, and known gaps

1. No real server has been inspected, so section 11 is not yet met. I need the cloud VM host and SSH access (open question 1).
2. The checker's parsing of Linux output (nvidia-smi, df, ss) is tested against sample text I wrote, not real output. Check it against the first real inspection.
3. The "df" check reads the "/" line in Linux format only; macOS output is not parsed.
4. scripts/test-all.sh calls "python", which does not exist on this Mac (only python3), so the full suite did not run here. I did not change it, and I did not add the new tests to it either.
5. The two plan approvals still needed: D2 to D5, and section 9 question 1.

Suggested next step

Provision the temporary GPU VM, run the inspection, then start Phase 1 (infra/inference) on it.